

Use of the library **fitdistrplus** to specify a distribution from non-censored or censored data

Marie Laure Delignette-Muller, Régis Pouillot and Jean-Baptiste Denis

November 28, 2008

Here you will find some easy examples of use of the functions of the library **fitdistrplus**. The aim is to show you by examples how to use these functions to help you to specify a parametric distribution from data corresponding to a random sample drawn from a theoretical distribution that you want to describe. For details, see the documentation of each function, using the R help command (ex.: `?fitdistr`). Do not forget to load the library using the function `library` before testing following examples.

```
> library(fitdistrplus)
```

Contents



1	Specification of a distribution from non-censored continuous data	2
1.1	Graphical display of the observed distribution	2
1.2	Characterization of the observed distribution	2
1.3	Fitting of a distribution	4
1.4	Simulation of the uncertainty by bootstrap	12
2	Specification of a distribution from non-censored discrete data	13
3	Specification of a distribution from censored data	17
3.1	Graphical display of the observed distribution	17
3.2	Fitting of a distribution	20

1 Specification of a distribution from non-censored continuous data

1.1 Graphical display of the observed distribution

First of all, the observed distribution may be plotted using the function `plotdist`.

```
> x1 <- c(6.4, 13.3, 4.1, 1.3, 14.1, 10.6, 9.9, 9.6, 15.3, 22.1,  
+       13.4, 13.2, 8.4, 6.3, 8.9, 5.2, 10.9, 14.4)  
> plotdist(x1)
```




1.2 Characterization of the observed distribution

Descriptive parameters of the empirical distribution may be computed using the function `descdist`. This function will also provide by default a skewness-kurtosis plot which may help you to select which distribution(s) to fit among the potential candidates.

```
> descdist(x1)  
  
summary statistics  
-----  
min:  1.3   max:  22.1  
median: 10.2  
mean:  10.4  
sample sd:  4.75  
sample skewness: 0.314  
sample kurtosis: 3.25
```

Cullen and Frey graph



In order to take into account the uncertainty of the estimated values of kurtosis and skewness, the data set may be bootstrapped by fixing the argument `boot` to an integer above 10 in `descdist`. `boot` values of skewness and kurtosis corresponding to the boot nonparametric bootstrap samples are then computed and reported in blue color on the skewness-kurtosis plot.


```
> descdist(x1, boot = 1000)
```

```
summary statistics
```

```
-----
```

```
min:  1.3   max: 22.1
median: 10.2
mean:  10.4
sample sd: 4.75
sample skewness: 0.314
sample kurtosis: 3.25
```

Cullen and Frey graph



1.3 Fitting of a distribution

One or more parametric distributions may then be fitted to the data set, one at a time, using the function `fitdist`. This function uses the maximum likelihood method if the argument `method="mle"` (or if it is omitted) or the matching moments method if the argument `method="mom"`. When fitting continuous distributions, Kolmogorov-Smirnov and Anderson-Darling statistics are computed and corresponding tests are performed when possible. Even if less appropriate for continuous distributions, the Chi-squared statistic is also computed when possible. For this calculation, cells are defined by the argument `chisqbreaks` or automatically defined from the data set and from the argument `meancount` (the approximate mean count per cell) which is fixed to $(4n)^{2/5}$ if omitted (with n the length of the data set). For more details, see the help of the function `fitdist`. Four goodness of fit plots are also provided.

Below is the result of a fit of a gamma distribution by maximum likelihood.

```
> f1g <- fitdist(x1, "gamma")
> plot(f1g)
> summary(f1g)
```

```
FITTING OF THE DISTRIBUTION ' gamma ' BY MAXIMUM LIKELIHOOD
PARAMETERS
```

	estimate	Std. Error
shape	3.575	1.140
rate	0.343	0.118

Loglikelihood: -54.4

```
GOODNESS-OF-FIT STATISTICS
```

```
----- Chi-squared -----
Chi-squared statistic: 7.93
Degree of freedom of the Chi-squared distribution: 3
Chi-squared p-value: 0.0475
```


!!! the p-value may be wrong
 with some theoretical counts < 5 !!!

!!! For continuous distributions, Kolmogorov-Smirnov and
 Anderson-Darling statistics should be preferred !!!


```
----- Kolmogorov-Smirnov-----
Kolmogorov-Smirnov statistic: 0.138
Kolmogorov-Smirnov test: not calculated
```

```
----- Anderson-Darling-----
Anderson-Darling statistic: 0.457
Anderson-Darling test: not rejected
```


Empirical and theoretical distr.




QQ-plot



Empirical and theoretical CDFs



PP-plot



Below is the result of another fit of the same distribution by matching moments.

```
> flgbis <- fitdist(x1, "gamma", method = "mom")
> summary(flgbis)
```

```
FITTING OF THE DISTRIBUTION ' gamma ' BY MATCHING MOMENTS
PARAMETERS
```

```
      estimate
shape    4.810
rate     0.462
```

```
-----
```

```
GOODNESS-OF-FIT STATISTICS
```

```
----- Chi-squared-----
```

```
Chi-squared statistic: 7.27
```

```
Degree of freedom of the Chi-squared distribution: 3
```

```
Chi-squared p-value: 0.0637
!!! the p-value may be wrong
      with some theoretical counts < 5 !!!
```

```
!!! For continuous distributions, Kolmogorov-Smirnov and
      Anderson-Darling statistics should be preferred !!!
```

```
----- Kolmogorov-Smirnov-----
Kolmogorov-Smirnov statistic: 0.144
Kolmogorov-Smirnov test: not calculated
```

```
----- Anderson-Darling-----
Anderson-Darling statistic: 0.471
Anderson-Darling test: not rejected
```

As can be seen in this returned summary, the automatic definition of the cells required to calculate the Chi-squared statistic does not give theoretical counts large enough to validate the use of the test in this example. It is often the case for small data sets. The observed and theoretical counts may be printed as below :

```
> flg$chisqtable

      obscounts theocounts
<= 5.2 3.0000000 2.8950753
<= 8.4 3.0000000 4.5964712
<= 9.9 3.0000000 2.1080076
<= 13.2 3.0000000 3.7064776
<= 14.1 3.0000000 0.7577383
> 14.1 3.0000000 3.9362300
```

Below is the fit of a lognormal distribution.

```
> f11 <- fitdist(x1, "lnorm")
> plot(f11)
> summary(f11)
```

```
FITTING OF THE DISTRIBUTION ' lnorm ' BY MAXIMUM LIKELIHOOD
PARAMETERS
```

	estimate	Std. Error
meanlog	2.197	0.147
sdlog	0.622	0.104

Loglikelihood: -56.5





```
-----
GOODNESS-OF-FIT STATISTICS
```

```
----- Chi-squared-----
Chi-squared statistic: 11.1
Degree of freedom of the Chi-squared distribution: 3
Chi-squared p-value: 0.0110
!!! the p-value may be wrong
      with some theoretical counts < 5 !!!
```

```
!!! For continuous distributions, Kolmogorov-Smirnov and
      Anderson-Darling statistics should be preferred !!!
```

```
----- Kolmogorov-Smirnov-----
Kolmogorov-Smirnov statistic: 0.178
Kolmogorov-Smirnov test: not calculated
```

```
----- Anderson-Darling-----
Anderson-Darling statistic: 0.793
Anderson-Darling test: rejected
```

Empirical and theoretical distr.**QQ-plot****Empirical and theoretical CDFs****PP-plot**

Below is the fit of a normal distribution.

```
> f1n <- fitdist(x1, "norm")
> plot(f1n)
> summary(f1n)
```

FITTING OF THE DISTRIBUTION ' norm ' BY MAXIMUM LIKELIHOOD
PARAMETERS

	estimate	Std. Error
mean	10.41	1.119
sd	4.75	0.791

Loglikelihood: -53.6

GOODNESS-OF-FIT STATISTICS

----- Chi-squared -----

Chi-squared statistic: 4.83

Degree of freedom of the Chi-squared distribution: 3

Chi-squared p-value: 0.185

!!! the p-value may be wrong

with some theoretical counts < 5 !!!

!!! For continuous distributions, Kolmogorov-Smirnov and
Anderson-Darling statistics should be preferred !!!

----- Kolmogorov-Smirnov -----


Kolmogorov-Smirnov statistic: 0.110

Kolmogorov-Smirnov test: not calculated


----- Anderson-Darling -----

Anderson-Darling statistic: 0.226
 Anderson-Darling test: not rejected


Empirical and theoretical distr.




QQ-plot



Empirical and theoretical CDFs



PP-plot



Below is the fit of a Weibull distribution.

```
> f1w <- fitdist(x1, "weibull")
> plot(f1w)
> summary(f1w)
```

FITTING OF THE DISTRIBUTION ' weibull ' BY MAXIMUM LIKELIHOOD
 PARAMETERS

	estimate	Std. Error
shape	2.29	0.426
scale	11.70	1.264

Loglikelihood: -53.5

GOODNESS-OF-FIT STATISTICS

----- Chi-squared -----
 Chi-squared statistic: 5.87
 Degree of freedom of the Chi-squared distribution: 3
 Chi-squared p-value: 0.118
 !!! the p-value may be wrong
 with some theoretical counts < 5 !!!


!!! For continuous distributions, Kolmogorov-Smirnov and
 Anderson-Darling statistics should be preferred !!!

----- Kolmogorov-Smirnov -----
 Kolmogorov-Smirnov statistic: 0.121


Kolmogorov-Smirnov test: not calculated

----- Anderson-Darling -----
Anderson-Darling statistic: 0.282
Anderson-Darling test: not rejected


Empirical and theoretical distr.




QQ-plot



Empirical and theoretical CDFs



PP-plot



The values of the Anderson-Darling statistic (or another result of the fit: see the help of `fitdist` for details) for the different fittings may be extracted and compared to help the selection of a distribution :

```
> anderson <- list(lnorm = f1l$ad, gamma = f1g$ad, norm = f1n$ad,  
+   weibull = f1w$ad)  
> anderson
```

```
$lnorm  
[1] 0.7925666
```

```
$gamma  
[1] 0.4567361
```

```
$norm  
[1] 0.225598
```


```
$weibull  
[1] 0.2821827
```

For some distributions (see the help of `fitdist` for details), it is necessary to specify initial values for the distribution parameters in the argument `start` when using the maximum likelihood method. `start` must be a named list of parameters initial values. The names of the parameters in `start` must correspond exactly to their definition in R or to their definition in a previous R code. The function `plotdist` may help to find correct initial values for the distribution parameters in non trivial cases, by an manual iterative use if necessary.


For example, below is the definition of the Gumbel distribution (also named extreme value distribution) and a first plot of the data set with the Gumbel distribution with arbitrary values for parameters.

```
> dgumbel <- function(x, a, b) 1/b * exp((a - x)/b) * exp(-exp((a -  
+ x)/b))  
> pgumbel <- function(q, a, b) exp(-exp((a - q)/b))  
> qgumbel <- function(p, a, b) a - b * log(-log(p))  
> plotdist(x1, "gumbel", para = list(a = 3, b = 2))
```


Empirical and theoretical distr.




QQ-plot



Empirical and theoretical CDFs







PP-plot



The same data set may be plotted with a Gumbel distribution with modified values for parameters.

```
> plotdist(x1, "gumbel", para = list(a = 10, b = 5))
```

Empirical and theoretical distr.**QQ-plot****Empirical and theoretical CDFs****PP-plot**

And a Gumbel distribution may be fitted to data with these values for initial parameter values.

```
> fgv <- fitdist(x1, "gumbel", start = list(a = 10, b = 5))
> plot(fgv)
> summary(fgv)
```

```
FITTING OF THE DISTRIBUTION ' gumbel ' BY MAXIMUM LIKELIHOOD
PARAMETERS
```

```
  estimate Std. Error
a      8.09      1.092
b      4.38      0.766
Loglikelihood: -54.1
```

```
-----
```

```
GOODNESS-OF-FIT STATISTICS
```

```
----- Chi-squared -----
```

```
Chi-squared statistic: 7.56
```

```
Degree of freedom of the Chi-squared distribution: 3
```

```
Chi-squared p-value: 0.056
```

```
!!! the p-value may be wrong
```

```
with some theoretical counts < 5 !!!
```

```
!!! For continuous distributions, Kolmogorov-Smirnov and
Anderson-Darling statistics should be preferred !!!
```

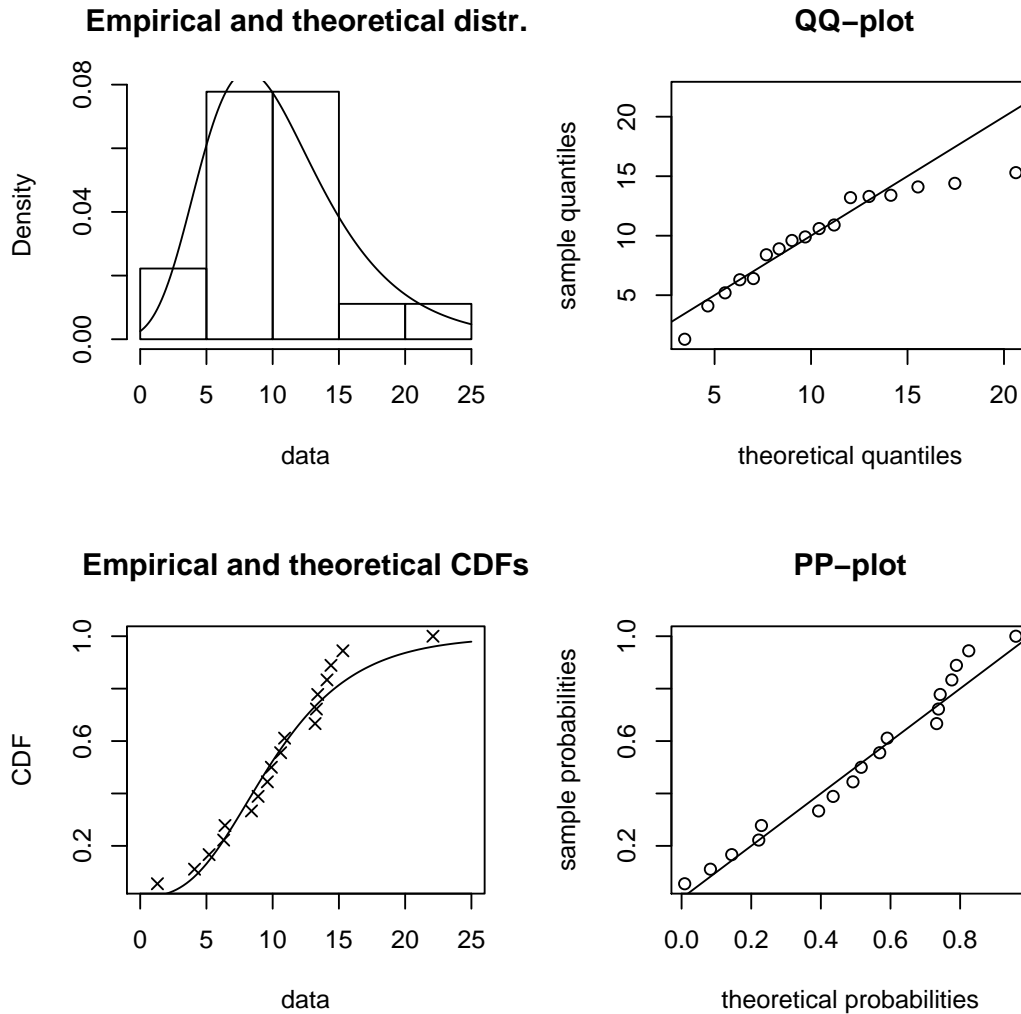
```
----- Kolmogorov-Smirnov -----
```

```
Kolmogorov-Smirnov statistic: 0.121
```

```
Kolmogorov-Smirnov test: not calculated
```

```
----- Anderson-Darling -----
```

Anderson-Darling statistic: 0.34
 Anderson-Darling test: not calculated



1.4 Simulation of the uncertainty by bootstrap

The uncertainty in the parameters of the fitted distribution may be simulated by parametric or nonparametric bootstrap using the function `bootdist`. This function returns the bootstrapped values of parameters which may be plotted to visualize the bootstrap region. It also calculates the 95 percent confidence intervals for each parameter from the 2.5 and 97.5 percentiles of the bootstrap values of each parameter (see the help of the function `bootdist` for details).

Below is an example of the use of this function with the previous fit of the gamma distribution.


```
> b1g <- bootdist(f1g)
> plot(b1g)
> summary(b1g)
```

Parametric bootstrap medians and 95% CI

	Median	2.5%	97.5%
shape	4.004	2.275	8.310
rate	0.386	0.212	0.824

Maximum likelihood method converged for 999 among 999 iterations

Scatterplot of boostrapped values of parameters




2 Specification of a distribution from non-censored discrete data


A discrete data set may be considered as a continuous one for example for a large data set from a binomial distribution converging to a normal one. A discrete plot of the distribution may also be provided, fixing the argument `discrete` of the function `plotdist` to `TRUE`.

```
> x2 <- rnbino(n = 100, size = 2, prob = 0.3)
> plotdist(x2, discrete = TRUE)
```

Empirical distribution



Empirical CDFs



As for continuous distributions, descriptive parameters of the empirical distribution may be computed using the function `descdist` which also provides a skewness-kurtosis plot which may help you to choose which distribution(s) to fit.

```
> descdist(x2, discrete = T)
```

```
summary statistics
```

```
-----
```

```
min: 0    max: 23
```

```
median: 3
```


```
mean: 4.35
```

```
sample sd: 4.17
```

```
sample skewness: 1.92
```

```
sample kurtosis: 7.95
```

Cullen and Frey graph



As for continuous distributions, one or more parametric distributions may then be fitted to the data set by maximum likelihood or matching moments.

Below is the result of the fit of a Poisson distribution with the bootstrap simulations.

```
> f2p <- fitdist(x2, "pois")
> plot(f2p)
> summary(f2p)
```

```
FITTING OF THE DISTRIBUTION ' pois ' BY MAXIMUM LIKELIHOOD
PARAMETERS
```

```
      estimate Std. Error
lambda      4.35      0.209
Loglikelihood:  -319
```

```
-----
GOODNESS-OF-FIT STATISTICS
```


```
----- Chi-squared -----
Chi-squared statistic:  141
Degree of freedom of the Chi-squared distribution:  6
Chi-squared p-value:  5.56e-28
!!! the p-value may be wrong
      with some theoretical counts < 5 !!!
```

```
> b2p <- bootdist(f2p)
> summary(b2p)
```


```
Parametric bootstrap medians and 95% CI
Median  2.5% 97.5%
      4.35 3.96 4.74
```

```
Maximum likelihood method converged for 999 among 999 iterations
```

Empirical (black) and theoretical (red) distr.



Empirical (black) and theoretical (red) CDFs



Below is the result of the fit of a negative binomial distribution with the bootstrap simulations.

```
> f2n <- fitdist(x2, "nbinom")
> plot(f2n)
> summary(f2n)
```

FITTING OF THE DISTRIBUTION ' nbinom ' BY MAXIMUM LIKELIHOOD
PARAMETERS

	estimate	Std. Error
size	1.61	0.329
mu	4.35	0.401

Loglikelihood: -255

GOODNESS-OF-FIT STATISTICS

----- Chi-squared -----

Chi-squared statistic: 3.88

Degree of freedom of the Chi-squared distribution: 5

Chi-squared p-value: 0.566


```
> b2n <- bootdist(f2n)
> summary(b2n)
```

Parametric bootstrap medians and 95% CI


	Median	2.5%	97.5%
size	1.63	1.14	2.54
mu	4.34	3.63	5.11

Maximum likelihood method converged for 999 among 999 iterations

Empirical (black) and theoretical (red) distr.



Empirical (black) and theoretical (red) CDFs



3 Specification of a distribution from censored data


Censored data may contain left censored, right censored and interval censored values, with several lower and upper bounds. Data must be coded into a dataframe with two columns, respectively named **left** and **right**, describing each observed value as an interval. The **left** column contains either NA for left censored observations, the left bound of the interval for interval censored observations, or the observed value for non-censored observations. The **right** column contains either NA for right censored observations, the right bound of the interval for interval censored observations, or the observed value for non-censored observations.

3.1 Graphical display of the observed distribution

First of all, the observed distribution may be plotted using the function `plotdistcens`. Data are reported directly as segments for interval, left and right censored data, and as points for non-censored data. For more details, see the help of the function `plotdistcens`.

```
> d1 <- data.frame(left = c(1.73, 1.51, 0.77, 1.96, 1.96, -1.4,
+   -1.4, NA, -0.11, 0.55, 0.41, 2.56, NA, -0.53, 0.63, -1.4,
+   -1.4, -1.4, NA, 0.13), right = c(1.73, 1.51, 0.77, 1.96,
+   1.96, 0, -0.7, -1.4, -0.11, 0.55, 0.41, 2.56, -1.4, -0.53,
+   0.63, 0, -0.7, NA, -1.4, 0.13))
> plotdistcens(d1)
```


Cumulative distribution plot



When left or right NA-values correspond to finite value (for example 0 for left NA-values of positive data), the arguments `leftNA` (or `rightNA`) must be affected to this finite value to ensure a correct plot of left (or right) censored observations, as in the example below.

```
> d2 <- data.frame(left = 10^(d1$left), right = 10^(d1$right))
> plotdistcens(d2, leftNA = 0)
```


Cumulative distribution plot



It is also possible to fix `rightNA` or `leftNA` to a realistic extreme value, even if not exactly known, to obtain a reasonable global ranking of observations, as in the example below for the first dataset.

```
> plotdistcens(d1, rightNA = 3)
```

Cumulative distribution plot



3.2 Fitting of a distribution

One or more parametric distributions may then be fitted to the censored data set, one at a time, using the function `fitdistcens`. This function always uses the maximum likelihood method. For more details, see the help of the function `fitdistcens`. Only one goodness of fit plot is provided for censored data, in cumulative frequencies. The uncertainty in the parameters of the fitted distribution may be simulated by nonparametric bootstrap only, using the function `boodistcens`.

Below is the result of a fit of a Weibull distribution by maximum likelihood and the results of the corresponding bootstrap simulations.


```
> f2w <- fitdistcens(d2, "weibull")
> summary(f2w)
```

```
FITTING OF THE DISTRIBUTION ' weibull ' BY MAXIMUM LIKELIHOOD ON CENSORED DATA
PARAMETERS
```

```
      estimate Std. Error
shape    0.324    0.0613
scale    6.124    4.5872
Loglikelihood: -68.5
```

```
> plot(f2w, leftNA = 0)
```

Cumulative distribution plot



```
> b2w <- bootdistcens(f2w)
> summary(b2w)
```


Nonparametric bootstrap medians and 95% CI

	Median	2.5%	97.5%
shape	0.338	0.245	0.481
scale	6.190	1.194	25.682

Maximum likelihood method converged for 999 among 999 iterations

```
> plot(b2w)
```

Scatterplot of the bootstrapped values of the two parameters



Goodness of fit statistics are not computed for fit on censored data, so the quality of fit may only be estimated from the loglikelihood and the goodness of fit plot.

Below is the fit of a lognormal distribution to the same censored data set.

```
> f2l <- fitdistcens(d2, "lnorm")
> summary(f2l)
```


FITTING OF THE DISTRIBUTION ' lnorm ' BY MAXIMUM LIKELIHOOD ON CENSORED DATA
PARAMETERS

	estimate	Std. Error
meanlog	0.27	0.764
sdlog	3.28	0.600

Loglikelihood: -68.7

```
> plot(f2l, leftNA = 0)
```

Cumulative distribution plot



Below is the fit of an exponential distribution.


```
> f2e <- fitdistcens(d2, "exp")
> summary(f2e)
```

FITTING OF THE DISTRIBUTION ' exp ' BY MAXIMUM LIKELIHOOD ON CENSORED DATA
PARAMETERS

	estimate	Std. Error
rate	0.0292	0.00668
Loglikelihood:	-99.6	

```
> plot(f2e, leftNA = 0)
```

Cumulative distribution plot



As with `fitdist`, for some distributions (see the help of `fitdistcens` for details), it is necessary to specify initial values for the distribution parameters in the argument `start`. `start` must be a named list of parameters initial values. The names of the parameters in `start` must correspond exactly to their definition in R or to their definition in a previous R code. The function `plotdistcens` may help to find correct initial values for the distribution parameters in non trivial cases, by an manual iterative use if necessary, as explained previously for non-censored continuous data.